



# Deep Learning with Differential Privacy

Kaoutar BOULIF, Charlotte  
DURAND



# Introduction

Article : « Deep Learning with Differential Privacy »

Authors : Martín Abadin and colleagues

Submission year : 2016

Contribution : combine machine learning algorithms with advanced privacy preserving mechanisms , training neural networks within a modest privacy budget

Implementation : Tensorflow

Datasets: MNIST / CIFAR-10

# Basic concepts

## DIFFERENTIAL PRIVACY:

*Definition 1.* A randomized mechanism  $\mathcal{M}: \mathcal{D} \rightarrow \mathcal{R}$  with domain  $\mathcal{D}$  and range  $\mathcal{R}$  satisfies  $(\varepsilon, \delta)$ -differential privacy if for any two adjacent inputs  $d, d' \in \mathcal{D}$  and for any subset of outputs  $S \subseteq \mathcal{R}$  it holds that

$$\Pr[\mathcal{M}(d) \in S] \leq e^\varepsilon \Pr[\mathcal{M}(d') \in S] + \delta.$$

## SENSITIVITY:

$$\Delta f = \max \|f(D_1) - f(D_2)\|_1$$

on datasets  $D_1, D_2$  differing on at most one element

# Differentially private SGD algorithm

---

**Algorithm 1** Differentially private SGD (Outline)

---

**Input:** Examples  $\{x_1, \dots, x_N\}$ , loss function  $\mathcal{L}(\theta) = \frac{1}{N} \sum_i \mathcal{L}(\theta, x_i)$ . Parameters: learning rate  $\eta_t$ , noise scale  $\sigma$ , group size  $L$ , gradient norm bound  $C$ .

**Initialize**  $\theta_0$  randomly

**for**  $t \in [T]$  **do**

    Take a random sample  $L_t$  with sampling probability  $L/N$

**Compute gradient**

    For each  $i \in L_t$ , compute  $\mathbf{g}_t(x_i) \leftarrow \nabla_{\theta_t} \mathcal{L}(\theta_t, x_i)$

**Clip gradient**

$\tilde{\mathbf{g}}_t(x_i) \leftarrow \mathbf{g}_t(x_i) / \max(1, \frac{\|\mathbf{g}_t(x_i)\|_2}{C})$

**Add noise**

$\tilde{\mathbf{g}}_t \leftarrow \frac{1}{L} (\sum_i \tilde{\mathbf{g}}_t(x_i) + \mathcal{N}(0, \sigma^2 C^2 \mathbf{I}))$

**Descent**

$\theta_{t+1} \leftarrow \theta_t - \eta_t \tilde{\mathbf{g}}_t$

**Output**  $\theta_T$  and compute the overall privacy cost  $(\varepsilon, \delta)$  using a privacy accounting method.

---

# Hyperparameter tuning and results

## MNIST:


Baseline model: 98.3% in about 100 epochs

### Differentially private model:


small noise scale ( $\sigma = 2$ ,  $\sigma_p = 4$ ): 90% accuracy for ( $\epsilon=0.5$ ,  $\delta=10^{-5}$ )DP

medium ( $\sigma = 4$ ,  $\sigma_p = 7$ ): 95% accuracy for ( $\epsilon=2$ ,  $\delta=10^{-5}$ ) DP


large ( $\sigma = 8$ ,  $\sigma_p = 16$ ): 97% accuracy for ( $\epsilon=8$ ,  $\delta=10^{-5}$ ) DP



(1) Large noise



(2) Medium noise



(3) Small noise


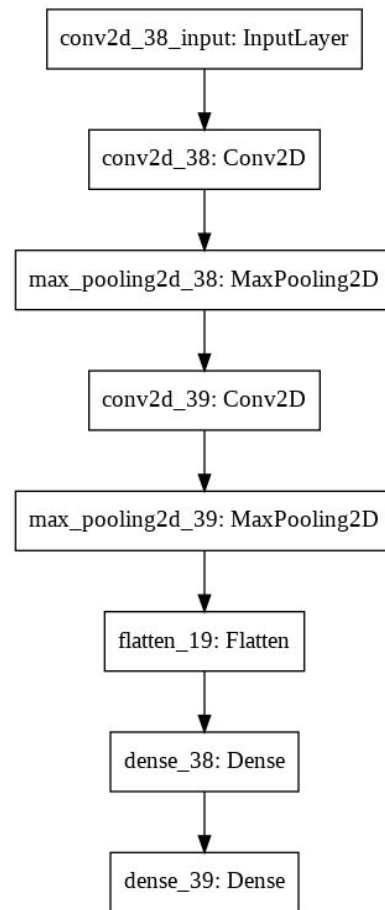
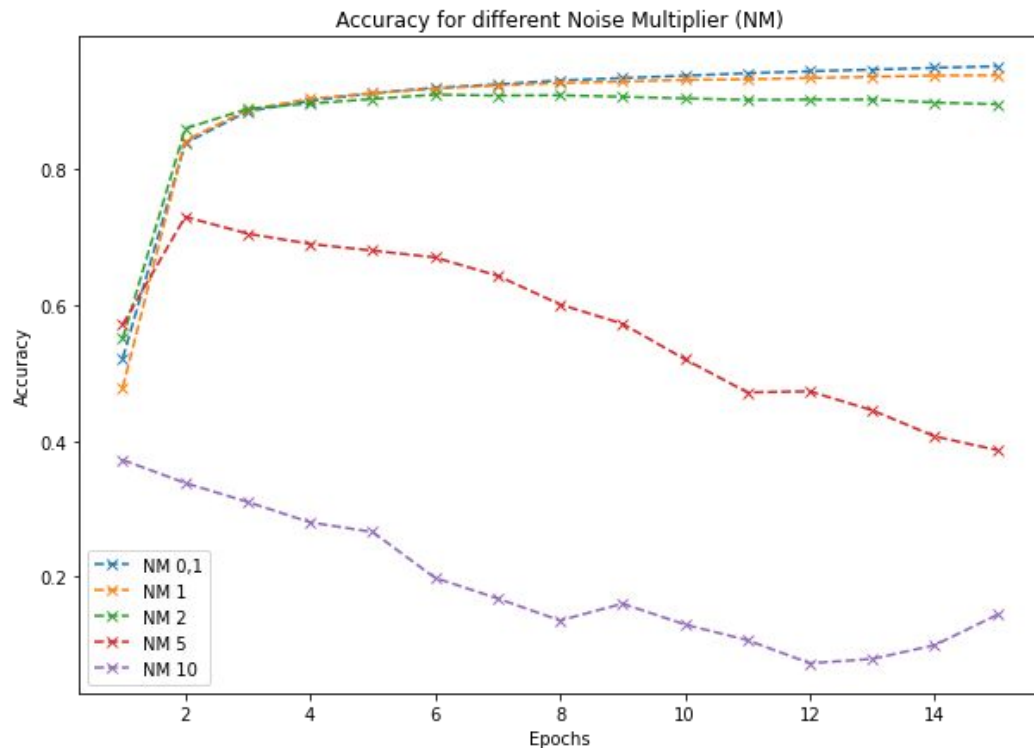


Figure 4: Accuracy of various  $(\epsilon, \delta)$  privacy values on the MNIST dataset. Each curve corresponds to a different  $\delta$  value.


# MNIST




Architecture du modèle

# MNIST

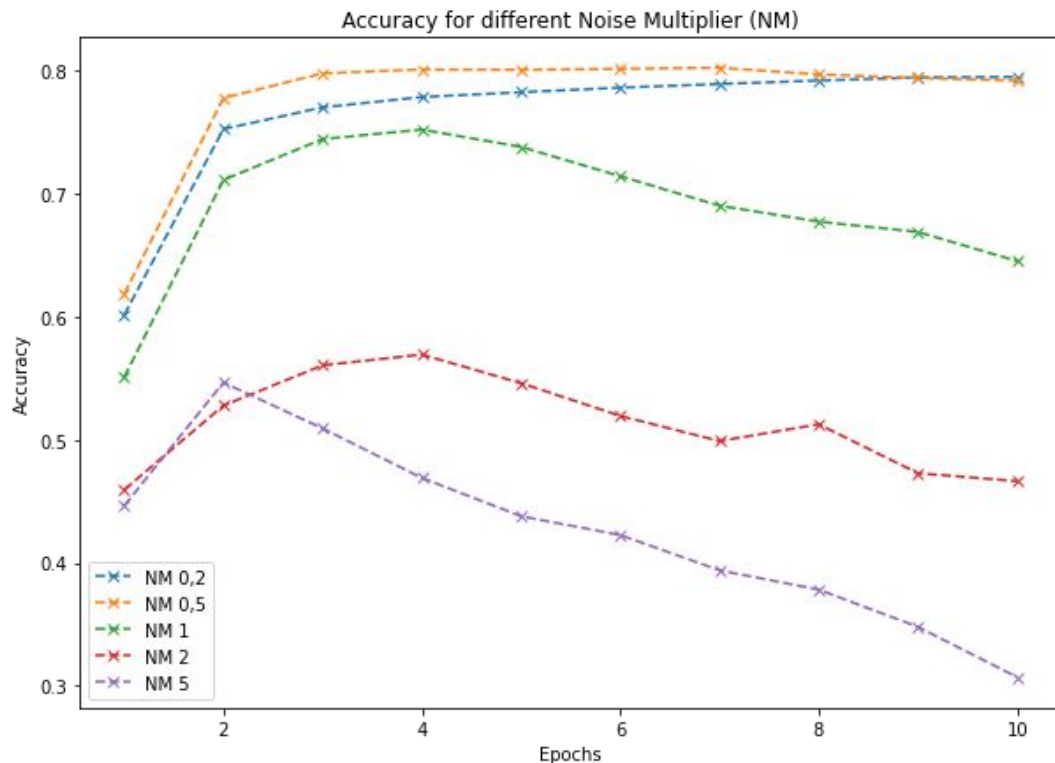
Evolution of the accuracy after 15 epoch depending on epsilon



evolution of epsilon for NM=1



# Fashion MNIST




- lower accuracy
- fine tuning hyperparameters :  
l2\_norm=5  
increase of the learning rate
- Longer computation
- Need to decrease the learning rate when Noise Multiplier is bigger



# Fashion MNIST

Evolution of the accuracy after 10 epoch depending on epsilon



Similar dependance than for MNIST

Find the good trade-off between accuracy and differential privacy

# Conclusion

→ Fine-tuning hyperparameters isn't easy

→ Yet, we obtain interesting results but computation is time consuming

Future possibilities :

- training with another optimizer (RMSProp for example)
- training on other dataset
- seeing the influence of the norm clipping

Thank you for your attention !

Any questions ?

